

CRDTs

Conflict-free Replicated Data Types

Mohammad Ghasemisharif

John Kristoff

Sepideh Roghanchi

Agenda

- Background and problem definition
- Methods of synchronization
 - Operation-based consistency
 - State-based consistency
- Algorithms, pseudo-code, and examples
 - Data types (graph, set, counter)
- Real-world usage and limitations

Simple Client/Server

Availability? Scale?



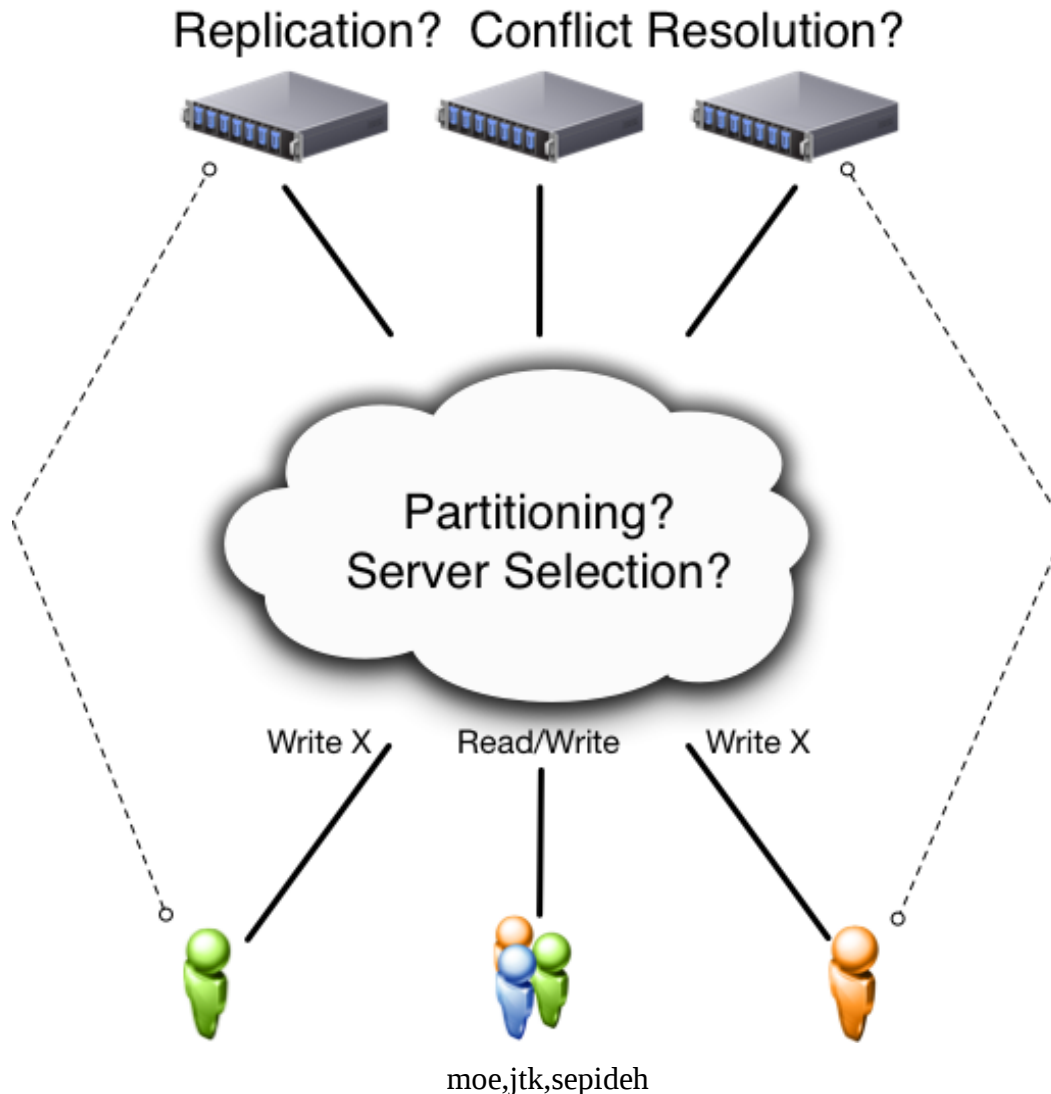
Read

Read/Write

Write



Distributed System



Conflicts and Resolution

- Partition A receives update(X_a)
- Partition B receives update(X_b)
- How to merge competing updates?
 - Timestamp?
 - Last-writer-wins?
 - Solution may be imperfect (i.e. a loser)
- Some data-types are conflict-free!

Replication

- Availability is desirable
- Fault-tolerance is desirable
- Low latency (high-performance) is desirable
- Maintaining consistency can be difficult

Key Concepts to Review

- Strong consistency
- Eventual consistency
- Conflict arbitration
 - Consensus and rollback
- CAP problem (pick 2)
 - Consistency
 - Availability
 - Partition Tolerance

strong consistency

- Everyone knows about every update immediately
- There is an order for all operations
- Everyone sees the same order
- Bottlenecks:
 - Consensus problem
 - Makes the system sequential
 - Slow, not scalable
 - Tolerates $< n/2$ failures

- Unfortunately, the CAP theorem tells us that no system satisfying those three desirable properties together exists:
 - Availability is often dropped!
- Drop strong consistency for a weaker form of consistency:
 - Eventual consistency


eventual consistency

- Update local + propagate
- All updates eventually take effect at all replicas, asynchronously and possibly in different orders
- Concurrent updates may conflict
- Still needs consensus
 - Conflict -> reconcile
 - Moved consensus off the critical path (background)
- Better performance, more complex
- May come at the cost of availability when synchronizing!

Motivation

- provide a data structure distributed over a large network and manipulated by a large base of users around the world
- Having multiple replicas of the data structure:
 - good for fault tolerance and read latency
 - Problem with updates :
 - Synchronize -> slow
 - Don't synchronize -> conflicts
- The provided data structure should follow the CAP properties:
 - Consistency
 - Availability
 - partition-tolerance

Strong eventual consistency

- Update local + propagate
- A replica of the shared data structure is coherent with other replicas that *have observed* the same operations
- No synchronization (no consensus)
- Deterministic outcome for every conflict
- Allow any number of failure
-  solves CAP problem

Origin Story

- 2011 publication by Marc Shapiro, et al.
- Strong consistency does not scale
- Eventual consistency is better, but...
- Concurrent conflict resolution is hard
- Therefore, strong eventual consistency

How to do it? Need data types to support it...

- CRDT: a simple, theoretically sound approach to eventual consistency
- replicas of any CRDT converge to a common state that is equivalent to some correct sequential execution
- Properties:
 - no synchronization
 - update executes immediately
 - unaffected by network latency, faults, or disconnection
 - extremely scalable
 - fault-tolerant
 - does not require much mechanism

CRDT

- Conflict-free
- Replicated
- Data Type

Data-type (in CRDTs)

- Data structures that ease consistency
- Eliminates complexity of consensus
- Data or operations must be commutative

Definitions

- EC:
 - **Eventual delivery:** *An update delivered at some correct replica is eventually delivered to all correct replicas: $\forall i,j : f \in c_i \Rightarrow \blacklozenge f \in c_j$*
 - **Termination:** *All method executions terminate*
 - **Convergence:** *Correct replicas that have delivered the same updates eventually reach equivalent state: $\forall i,j : c_i = c_j \Rightarrow \blacklozenge s_i \equiv s_j$*
- SCE
 - **Strong Convergence:** *Correct replicas that have delivered the same updates **have** equivalent state: $\forall i,j : c_i = c_j \Rightarrow s_i \equiv s_j$*

Op versus State-based

- Operation: add 5, subtract 6
- State: $\text{send}(x_i)$, $\text{send}(x_{i+1})$

Commutative

- Data that commutes is a key property
 - e.g. $\text{add } A, \text{add } B == \text{add } B, \text{add } A$
- Not all ops or states are commutative
 - e.g. $\text{multiple } 5, \text{subtract } 6 \neq \text{subtract } 6, \text{multiply } 5$

Operation-based replication

- Local replica sends operation
- Concurrent operations must be commutative

op-based object

- Tuple: (S, s^0, q, t, u, P)
- S, s^0 and q same as state-based object
- Replica at process p_i has state $s_i \in S$
- Initial state is s^0
- $t =$ side-effect-free *prepare-update*
- $u =$ *effect-update*
- $P =$ delivery precondition


State-based replication

- Local replica (p1)
 - performs computation
 - updates local state
 - Periodically send(p1_state)
- Receiver replica(s)
 - perform merge(p1_state, p2_state)

state-based object


- Tuple: (S, s^0, q, u, m)
- Replica at process p_i has state $s_i \in S$
- Initial state is s^0
- q = query
- u = update
- m = merge

- Clients send query to read replica's state
 - Read only -> easy
- Updates:
 - State-based
 - Operation-based



State-based approach


- State-based:
 - Local queries, local updates at source
 - Send full state every once in a while
 - On receive, merge
 - File systems (NFS, Unison, Dynamo)




- all the replicas talk to each other directly or indirectly
- if it converges it will satisfy the strong consistency
- What is the sufficient condition for it to converge?

- Semilattice
 - Set with partial order and an operation can take any two values and give you an upper bound on them
- If the payload forms a semilattice (partial order on values, always can take an upper bound)
- If updates are increasing in semilattice
- If merge function computes this upper bound
- -> replicas converge to LUB of last values

State-based example




Another state-based example




- Out of order delivery

Operation-based approach

- operation-based:
 - Only updates are sent (smaller)
 - Each replica replay the updates
 - Reconcile non-commutative operations
 - Collaborative editing, Bayou, PNUTS
- Need something stronger!
- Make sure that all updates are propagated to all other replicas (downstream replicas)



Op-based example



$$(5 + 4) \times 2 \neq (5 \times 2) + 4$$

- Updates should be commutative operations

Compare!

- state-based:
 - Update and merge
 - Simple data types
 - Not efficient for large objects
- Operation-based:
 - Update
 - More complex
 - More powerful
 - Small messages
- They are equivalent
 - you can take any state based object and emulate it in an op-based model and if one converges the other converges

G-Counter CRDT

- TODO: algorithm and method (see wikipedia page)

PN-Counter CRDT

- TODO: algorithm and method (see wikipedia page)

G-Set CRDT

- TODO: algorithm and method (see wikipedia page)

2P-Set CRDT

- TODO: algorithm and method (see wikipedia page)

LWW-Set CRDT

- TODO: algorithm and method (see wikipedia page)

OR-Set CRDT

- TODO: algorithm and method (see wikipedia page)

Sequence CRDTs

- TODO: algorithm and method (see wikipedia page)

CRDT in Action

- SoundCloud
- Bet365
- Redis
- Riak
- League of Legions
- orbit-db

References

- <https://pages.lip6.fr/Marc.Shapiro/>
- <https://www.youtube.com/watch?v=ebWVLVhiaiY>
- <https://www.youtube.com/watch?v=vBU70EjwGfw>
- https://en.wikipedia.org/wiki/Conflict-free_replicated_data_type